
CHES-WORD-MODEL: A 10M-Game Benchmark for Exact State Tracking from Chess Move Sequences

Benjamin Walker

Mathematical Institute, University of Oxford
contact@benwalker.co.uk

Terry Lyons

Mathematical Institute, University of Oxford
Department of Mathematics, Imperial College London

Abstract

World models require state tracking, which is the ability to maintain a correct latent state across action sequences. Existing benchmarks are often synthetic or language-based, limiting their value as tests of structured state updates in realistic domains. We introduce CHES-WORD-MODEL, a large-scale state-tracking benchmark built from 10 million real chess games, where models predict the exact board state reached after a sequence of legal moves. Alongside a held-out real-game split, we include an out-of-distribution split from uniformly random legal play, which tests whether models learn the transition rules rather than shortcuts from common human positions. Prior theoretical and empirical work has shown that Transformers struggle to state-track, while input-dependent linear RNNs require expressive state-transition matrices to do so. We therefore benchmark a causal Transformer, block-diagonal SLiCE, Mamba-3, and Gated DeltaNet with negative eigenvalues under a matched interface and training protocol. The recurrent models strongly outperform the Transformer at 3 and 8 million parameters. Real-game performance saturates above 18 million parameters, but the random-uniform split remains discriminative up to 38 million, exposing failures otherwise hidden by scale. Additionally, ablations show that less expressive state-transition mechanisms reduce performance on the out-of-distribution split for all three recurrent models. Together, these results establish CHES-WORD-MODEL as a practical large-scale benchmark for state tracking that exposes failures model scale would otherwise conceal.

1 Introduction

1.1 Motivation

A world model is only as useful as the state it maintains. As actions and observations accumulate, the model must update this state, preserve information over long horizons, and remain consistent when trajectories depart from familiar training patterns. This requirement appears in classical model-based reinforcement learning and predictive state representations (Sutton, 1990; Littman et al., 2001), in neural world models that learn latent dynamics for control (Ha et al., 2018; Hafner et al., 2020), and in recent joint-embedding methods that learn representations by predicting images and videos in latent space (LeCun, 2022; Assran et al., 2023; Assran et al., 2025). In all of these settings, prediction is not merely a one-step pattern-matching problem. It depends on whether the model can maintain and update a coherent representation of the world.

Despite this central role, exact state maintenance remains difficult to evaluate. Many world-model tasks couple state tracking to perception, reward prediction, planning, or language understanding, making failures hard to localise. Conversely, many formal state-tracking benchmarks isolate the computation cleanly but rely on synthetic tasks such as parity, modular arithmetic, regular languages, or permutation composition (Delétang et al., 2023; Cirone et al., 2024; Grazzi et al., 2025; Merrill et al., 2024). These tasks have been crucial for exposing architectural limitations. However, because they are often small or highly regular, increasing model scale can obscure state-tracking failures by allowing models to exploit task-specific shortcuts or memorise large parts of the problem structure (Delétang et al., 2023; Li et al., 2025). They therefore leave open a complementary empirical question. Can modern sequence models maintain a structured latent state in a realistic deterministic action domain?

Chess provides a natural testbed for this question. The task is not to play chess, but to act as a learned game engine. Given only a sequence of legal moves, the model must reconstruct the exact state reached by the game. This state includes both the board configuration and the auxiliary variables needed to specify the position. The transition rules are known exactly, and every prediction can be checked against the rules of chess without learned simulators or subjective judgments.

Chess also provides a principled axis of distribution shift. Human games occupy a structured and strategically meaningful subset of legal play, just as observations of the physical world occupy only a small subset of all dynamically possible trajectories. However, a world model should not merely interpolate within familiar trajectories. It should infer the underlying dynamics well enough to remain consistent under distribution change. Uniformly random legal play tests this distinction directly. It produces positions that are valid under the same transition rules, but that depart sharply from human regularities. Strong performance on held-out human games may therefore reflect memorised openings, common tactical motifs, or likely piece configurations. Strong performance under random legal play requires robust rule-consistent state updates.

1.2 Contributions

We introduce **CHESS-WORLD-MODEL**, a large-scale benchmark for exact state tracking from chess move sequences. The benchmark is constructed from 10M real games and evaluates models on both held-out real games and a random-uniform split generated by uniformly random legal play. Given a sequence of moves, models must predict the corresponding sequence of game states, including piece locations and auxiliary state variables. This sequence-to-sequence formulation allows every prediction to be evaluated exactly, while the random-uniform split tests whether models have learned the transition rules of chess rather than shortcuts from the distribution of human games.

Recent theoretical and empirical work has shown that Transformers and input-dependent linear RNNs with diagonal state-transition matrices, such as Mamba and Mamba-2, struggle with state-tracking tasks (Merrill et al., 2022; Merrill et al., 2024; Walker et al., 2025). By contrast, richer recurrent transition mechanisms, such as input-dependent state-transition matrices with negative eigenvalues, block-diagonal structure, or diagonal-plus-low-rank structure, are theoretically more expressive and empirically improve state-tracking performance (Grazzi et al., 2025; Fan et al., 2024; Walker et al., 2025; Siems et al., 2025; Lahoti et al., 2026).

Motivated by these results, we benchmark a causal Transformer against three recent input-dependent linear RNNs with expressive state transitions: block-diagonal SLiCE, Mamba-3, and Gated DeltaNet with negative eigenvalues (Walker et al., 2025; Grazzi et al., 2025; Lahoti et al., 2026). All models are trained under a shared prediction interface and matched training protocol across four architecture scales, d128/l1, d256/l2, d384/l4, and d512/l6, where d denotes hidden dimension and l denotes number of layers. These correspond to approximately 3, 8, 18, and 38 million parameters, respectively.

Our experiments yield three main findings. First, the recurrent models substantially outperform the Transformer at the two smaller scales, d128/l1 and d256/l2. Second, held-out real-game performance largely saturates by d384/l4, whereas the random-uniform split remains discriminative through the largest scale, d512/l6. Third, ablations show that reducing the expressivity of the state-transition mechanism reduces random-uniform performance among the recurrent models. Together, these results establish **CHESS-WORLD-MODEL** as a practical benchmark for state tracking that exposes failures hidden by scale and in-distribution evaluation. It provides a foundation for developing

expressive and efficient sequence models whose update mechanisms support reliable state tracking, a necessary ingredient for learned world models.

1.3 Related work

Chess is a natural domain for studying state tracking in sequence models, since move sequences induce an exactly defined latent board state. Toshniwal et al. (2022) introduced chess as a testbed for language-model state tracking, training Transformer language models to predict move sequences and evaluating legal move prediction together with probes of the internal board state. Related interpretability work has studied whether chess models trained on game transcripts learn internal representations of board state and player skill (Karvonen, 2024). Cooper (2025) introduced the PGN-to-FEN benchmark, which tests whether pretrained large language models can reconstruct the exact final board state from a move sequence. Harang et al. (2025) introduced a closely related state-affordance evaluation that also prompts language models to infer chess states, but scores predictions by comparing the legal continuations induced by the predicted and true positions. On the theoretical side, Merrill et al. (2024) formalised chess in source-target move notation as an example of state tracking.

More broadly, deterministic board games have been used to study whether sequence models learn implicit world models. In Othello, models trained on move sequences have been shown to contain internal representations aligned with the underlying board state (Li et al., 2023; Nanda et al., 2023). At the same time, recent work argues that strong predictive performance can coexist with incoherent or inconsistent recovered world states, motivating evaluations that directly test state recovery rather than only next-action prediction (Vafa et al., 2024). This motivates using chess not only as a language-modelling domain, but as an exactly checkable state-tracking environment.

CHESS-WORLD-MODEL builds on this line of work by turning chess state tracking into a large-scale supervised benchmark for sequence-model architectures. It closely instantiates the formal source-target chess state-tracking problem studied by Merrill et al. (2024), while extending the target state to include auxiliary FEN-style variables. It differs from prior chess-based empirical evaluations in three main respects. First, it trains models to predict the complete game state at every prefix of a move sequence, rather than using next-move training with state probes or prompting pretrained language models for a final position (Toshniwal et al., 2022; Cooper, 2025; Harang et al., 2025). Second, it evaluates exact sequence-to-sequence state reconstruction, rather than only final-position recovery or indirect consistency through legal continuations. Third, it includes an out-of-distribution random-agent split, generated by uniformly random legal play, which tests whether models learn rule-consistent state evolution beyond the distribution of human games. Together, these choices make CHESS-WORLD-MODEL suitable for comparing sequence-model architectures as candidate state-update mechanisms for learned world models.

2 CHESS-WORLD-MODEL

2.1 Chess as a deterministic state-tracking environment

Standard chess is a simple world. The initial state is fixed, the dynamics are deterministic, and the action space is fully observed, regular, and noiseless. Unlike domains with perceptual ambiguity, stochastic transitions, or hidden exogenous variables, chess allows us to study state tracking in a setting where uncertainty about the world comes primarily from the need to integrate an action history correctly. This makes the domain useful for isolating the memory and update mechanisms of sequence models from confounds introduced by sensing, control, or environment noise.

We use the standard rules of chess as specified by the FIDE Laws of Chess (FIDE, 2023), and encode moves using UCI notation (Meyer-Kahlen et al., 2000). As shown in Figure 1, board squares are indexed by standard algebraic coordinates. Columns, called files, are labelled $a, \dots, h$, and rows, called ranks, are labelled $1, \dots, 8$. A move is encoded by concatenating its source and target squares. For example, the move $d3c4$ denotes a piece moving from square $d3$ to square $c4$. Promotion moves are encoded by appending the promoted piece type, for example $e7e8q$ for promotion to a queen.

Although each move is directly observed, the game state is not provided to the model. It must be reconstructed from the move stream alone. This requires more than remembering the current piece placement. In this paper, we use game state to mean the full FEN-style chess state reached after a

a8	b8	c8	d8	e8	f8	g8	h8
a7	b7	c7	d7	e7	f7	g7	h7
a6	b6	c6	d6	e6	f6	g6	h6
a5	b5	c5	d5	e5	f5	g5	h5
a4	b4	c4	d4	e4	f4	g4	h4
a3	b3	c3	d3	e3	f3	g3	h3
a2	b2	c2	d2	e2	f2	g2	h2
a1	b1	c1	d1	e1	f1	g1	h1

(a) The standard algebraic coordinates for board squares. Columns, called files, are labelled $a, \dots, h$, and rows, called ranks, are labelled $1, \dots, 8$.



(b) A move in UCI notation is represented by concatenating its source and target squares. Here the highlighted move is d3c4.

Figure 1: Chess move notation used in our dataset. The left panel shows the board coordinates used to identify squares. The right panel illustrates UCI move notation (Meyer-Kahlen et al., 2000), in which each move is recorded by its source square followed by its target square.

move prefix (Edwards, 1994), which includes whose turn it is, which castling rights remain available, whether an en passant capture is currently legal, the halfmove clock, and the fullmove number, but excludes the history of previous board states needed to adjudicate repetition-based draws. These quantities are part of the represented state because they affect legal continuations, draw-related counters, or the formal description of the position. The complete decomposition of this represented state into categorical prediction targets is given in Appendix A.1.

2.2 Task and trajectory alignment

The benchmark is a move-to-state trajectory prediction task rather than a terminal-board reconstruction task. Each example is a complete legal chess game beginning from the standard initial position. The input is the move sequence, while the target is the complete chess state after each observed prefix of that sequence.

Figure 2 illustrates the alignment. The first input token is a start symbol, aligned with the initial state. Each subsequent input token is the categorical encoding of a UCI move, aligned with the state obtained after applying that move. Thus, if the observed plies are $m_{1:T}$, the data pipeline constructs an aligned sequence of move tokens $x_{0:T}$ and chess states $s_{0:T}$ such that

$$x_0 = \text{<START>}, \quad s_0 = s_{\text{init}},$$

and for each $t \geq 1$,

$$x_t = \text{enc}_{\text{move}}(m_t), \quad s_t = \text{State}(m_{1:t}).$$

Here enc_{move} is the fixed move encoding map described in Section 2.3: it maps each UCI move to a single categorical token by its source square, target square, and promotion type. The model receives the encoded move prefix $x_{0:t}$ and predicts the full chess state s_t at every timestep. Every game therefore contributes supervision at every observed prefix length, including the initial position. Appendix B.1 provides additional details on PGN replay, trajectory construction, filtering, and deterministic evaluation.

2.3 Move and state encoding

Moves are represented with a fixed packed vocabulary over $64 \times 64 \times 5 = 20,480$ move geometries. The first two factors encode the source and target squares, and the final factor encodes the promotion type. The promotion value is one of no promotion, queen, rook, bishop, or knight. The implementation




Figure 2: CHESSE-WORLD-MODEL as an aligned move-to-state prediction task. The input sequence consists of a start token followed by UCI moves. Each input prefix is aligned with the complete chess state obtained after applying that prefix. The target state contains the piece occupying each square, the side to move, castling rights, en passant information, and the halfmove and fullmove counters.

appends two special symbols, `<START>` and `<PAD>`, giving a total move vocabulary size of 20,482. This encoding does not attempt to enumerate only legal moves. Legality is enforced by the source trajectories. Appendix A.2 reports move-token coverage and verifies that every packed move id appearing in the random-uniform test set is observed in the training data.

Each target state is represented as 75 categorical labels. The first 64 labels encode the board in square order $a8, b8, \dots, h1$, using the piece-state vocabulary $\mathcal{P} = \{\text{empty}\} \cup \{\text{white } p, \text{black } p \mid p \in \{\text{pawn, knight, bishop, rook, queen, king}\}\}$. The remaining 11 labels encode the FEN-style auxiliary state reached after the move prefix (Edwards, 1994), excluding the history of previous board states needed to adjudicate repetition-based draws. These labels encode the side to move, castling rights, en passant availability, the halfmove clock, and the fullmove number. The full per-label layout is given in Appendix A.1, and the same grouping is shown schematically in Figure 2.

2.4 Exactness metrics

The main body of the paper focuses on exact state accuracy, since the benchmark is designed to test whether a model recovers the complete chess state induced by a move prefix. Let J denote the set of 75 state targets, and write $\hat{s}_{t,j}$ for the prediction of target j at timestep t . The timestep-level exact state metric is $\text{ExactState}(t) = \mathbf{1}\{\hat{s}_t = s_t\}$. This score is equal to 1 only when all 75 state labels are predicted correctly. A single error in piece placement, side to move, castling rights, en passant state, or the move counters makes the state incorrect. Appendix B.1 defines additional metrics for partial state accuracy, trajectory exactness, and length-binned prefix exactness, with the corresponding results reported in Appendices D, E, and F.

2.5 Data source and split construction

The real-game corpus is derived from the public Lichess open database, whose monthly PGN exports are released under a CC0 license (Lichess, 2026). During preprocessing, each retained game is converted into the aligned move-state trajectory described in Section 2.2. The processed examples also store lightweight bookkeeping fields, such as game identifier, result, and length. Usernames, clock comments, annotations, and other metadata that are unnecessary for the benchmark are not used as modelling inputs. Games shorter than 10 full moves are filtered out.

The current benchmark instantiation uses 10M real games as the training and validation source. Training and validation are game-disjoint. Rather than materialising a separate validation file, the implementation assigns each game to train or validation by hashing its game identifier into a fixed

Split	Size	Description
Real-game train and validation	10M source games	Processed Lichess games
Held-out real validation	$\approx 0.5\%$ by default	Deterministic game split used for model selection
Held-out real test	10k	Fixed real-game benchmark set
Random-uniform test	10k	Fixed legal self-play benchmark set

Table 1: Benchmark splits in the current implementation. The main distribution shift is from human play to uniformly random legal self-play, while keeping the same chess dynamics and target representation.

bucket set. With the default settings, approximately 0.5% of games are reserved for validation. This gives a deterministic split that is reproducible across runs.

Final evaluation uses two fixed 10k-game test conditions. The first is a held-out real-game split from the same overall source family. This measures generalisation within the distribution of human play. The second is generated by sampling uniformly from the legal move set at each ply until game termination, then retaining 10k games of length at least 10 full moves. This measures robustness under legal but non-human action distributions. The four data splits are summarised in Table 1. Further details on preprocessing, the hash-based validation split, random-uniform generation, and evaluation protocol are given in Appendix B.1.

3 Experimental methodology

3.1 Shared prediction interface

All architectures are evaluated through the same move-to-state prediction interface. Each model receives the packed move sequence $x_{0:T}$ and predicts the aligned chess state $s_{0:T}$ at every timestep. The output layer is factorised into target-specific categorical heads. It consists of 64 board heads with 13 classes each, one binary side-to-move head, four binary castling heads, one 9-way en passant file head, one 3-way en passant rank head, and two 256-way byte heads for each of the halfmove and fullmove counters.

For a trajectory $(x_{0:T}, s_{0:T})$, the training objective is the masked sum of per-target cross-entropies over all non-padded timesteps,

$$\mathcal{L}(x_{0:T}, s_{0:T}) = -\frac{1}{|J|(T+1)} \sum_{t=0}^T \sum_{j \in J} \log p_{\theta}(s_{t,j} \mid x_{0:t}),$$

where J denotes the set of 75 state targets. This keeps the supervision signal, output representation, and evaluation metrics fixed across model families. Differences in performance therefore reflect how each architecture processes the move history, rather than differences in the prediction task.

The main architectural distinction is between attention over the whole observed prefix and recurrent state updates. A causal Transformer can revisit earlier moves through self-attention at every timestep. The recurrent models instead compress the observed history into a hidden state that is updated online as new move tokens arrive. Chess is useful for this comparison because the input stream is simple, but the latent state that must be maintained is exact, structured, and history-dependent.

3.2 Architectures

We compare four sequence-model families. The first is a causal Transformer (Vaswani et al., 2017), which serves as an attention-based baseline. The remaining three are recurrent architectures with input-dependent state updates: block-diagonal SLiCE (Walker et al., 2025), Mamba-3 (Gu et al., 2024; Dao et al., 2024; Lahoti et al., 2026), and Gated DeltaNet (Yang et al., 2025). For Gated DeltaNet, the main configuration allows negative eigenvalues, following recent evidence that this improves state-tracking behaviour (Grazzi et al., 2025). These recurrent architectures were chosen because they have been shown to perform well on state-tracking tasks and each has a less expressive variant for ablation: diagonal SLiCE, Mamba-2, and Gated DeltaNet with positive eigenvalues, respectively.




Figure 3: Test performance versus parameter count. The left panels show cross-entropy loss, and the right panels show exact state accuracy. The top row reports held-out real-game performance, while the bottom row reports random-uniform performance. Points show individual model families, and the dashed curve shows the empirical Pareto frontier over the evaluated configurations.

Each architecture is instantiated along the same width-depth schedule,

$$d_{128}/l_1, \quad d_{256}/l_2, \quad d_{384}/l_4, \quad d_{512}/l_6,$$

where d denotes hidden dimension and l denotes the number of layers. For the Transformer and Gated DeltaNet, these four settings use 2, 4, 6 and 8 heads, respectively. Mamba-3 uses the MIMO variant with rank 4. The resulting models define a common capacity axis rather than exactly matched parameter counts, which range from 3.13M to 3.20M at d_{128}/l_1 , 7.70M to 8.25M at d_{256}/l_2 , 16.89M to 19.72M at d_{384}/l_4 , and 35.14M to 40.83M at d_{512}/l_6 . Exact parameter counts and additional implementation details for each architecture are provided in Appendix B.2.

3.3 Optimisation

Training uses a two-stage procedure. The first stage is a one-epoch hyperparameter sweep over architecture, scale, learning rate, weight decay, warmup length, and Adam β_2 . The sweep covers the four architectures and four matched scales described above. It uses batch size 128 and AdamW optimisation, with learning rates $\{10^{-4}, 3 \times 10^{-4}, 10^{-3}, 3 \times 10^{-3}\}$, weight decays $\{0, 10^{-3}, 10^{-2}\}$, warmup steps $\{5000, 10000, 20000\}$, and Adam β_2 values $\{0.95, 0.99, 0.999\}$. This gives 1728 one-epoch runs. The first-stage sweep uses linear warmup as the learning rate scheduler. All first-stage runs use the same seed. Validation during the sweep is performed on a fixed number of held-out real-game batches, rather than by running a full validation pass after every checkpoint. The best first-stage configuration for each family and size is reported in Appendix C.

The second stage retrain the best-performing first-stage configurations from scratch for four epochs. These training runs use linear warmup followed by cosine decay, and tune the final learning-rate floor as a fraction of the maximum learning rate, with values in $\{0.1, 0.01\}$. Appendix C reports the second stage results, cosine-decay floor comparison, and final selected optimisation hyperparameters. This design is chosen to allow for a broad optimisation hyperparameter selection using the first stage and longer training runs of the selected configurations in the second stage.

4 Results

Figure 3 shows test performance as model scale increases. On the held-out real-game split, all recurrent models improve rapidly with scale and reach near-perfect exact state accuracy by approx-




Figure 4: Effect of reducing the expressivity of the recurrent state-transition mechanism. Filled markers denote the main recurrent variants, while open markers denote less expressive variants. Within each shaded size group, horizontal offsets are used only to separate markers visually and do not encode an additional variable. Arrows connect each less expressive variant to its corresponding main variant at the same nominal size. The main variants are block-diagonal SLiCE, Mamba-3, and Gated DeltaNet with negative eigenvalues. The less expressive variants are diagonal SLiCE, Mamba-2, and Gated DeltaNet with positive eigenvalues.

imately 8M parameters. As highlighted by the Pareto frontier, the Transformer is markedly less parameter-efficient at small and medium scales, obtaining substantially higher test loss and lower exact state accuracy than the recurrent models. This suggests that chess state tracking is sensitive not only to model capacity, but also to the structure of the state-update mechanism. Architectures with explicit recurrent updates appear better matched to maintaining an evolving latent state from a stream of legal moves. However, by the largest scales, held-out real-game performance is close to saturated for all model families, and exact state accuracy no longer clearly separates the strongest recurrent configurations.

The random-uniform split gives a different picture. Performance continues to improve across the full scale range, and the task remains discriminative even when the held-out real-game split has largely saturated. The strongest model reaches around 80% exact state accuracy on random-uniform games, while several models that are essentially perfect on held-out real games remain substantially below this level under random legal play. This indicates that high in-distribution accuracy does not imply that the model has learned robust rule-consistent state updates. Instead, held-out real-game evaluation can conceal failures that become visible only when the action distribution departs from human play. The complete held-out real-game and random-uniform test metrics underlying Figure 3 are reported in Appendix D. Appendix F further breaks down random-uniform performance by prefix length, showing where exact-state errors emerge along the trajectories.

Figure 4 isolates the effect of recurrent transition expressivity. For each recurrent family, the less expressive variant is compared with the corresponding main variant at the same nominal size. On the held-out real-game split, the gap is most visible at smaller scales. As scale increases, the main variants and their less expressive counterparts all approach near-perfect exact state accuracy, again showing that in-distribution evaluation becomes weakly discriminative once models are sufficiently large.

The random-uniform split reveals a much clearer and more persistent effect. Across SLiCE, Mamba, and Gated DeltaNet, replacing the main transition mechanism with a less expressive variant consistently reduces exact state accuracy and increases test loss. The most extreme case is diagonal SLiCE, which exhibits numerical instability on some out-of-distribution examples and produces very large

random-uniform cross-entropy despite reasonable held-out real-game performance. More broadly, the degradation remains visible at large scale, where held-out real-game performance is already saturated. This indicates that transition expressivity is not merely a small-model effect. It continues to matter for robust state tracking under distribution shift, and may help models learn rule-consistent updates rather than shortcuts that perform well only in distribution. Full held-out real-game and random-uniform ablation results for these less expressive variants are given in Appendix E.

Overall, these results support three conclusions. First, CHESS-WORLD-MODEL separates architecture families at small and medium scale, with recurrent models outperforming the Transformer under the matched interface. Second, held-out real-game evaluation alone is insufficient, because performance saturates and can make distinct architectures appear similarly strong. Third, the random-uniform split exposes failures in state tracking that are hidden by in-distribution evaluation, and shows that more expressive recurrent transition mechanisms give more robust rule-consistent state updates.

5 Limitations and societal impact

CHESS-WORLD-MODEL isolates exact state tracking in a deterministic action domain, but it does not capture all difficulties of world modelling. Chess removes perceptual uncertainty, stochastic transitions, partial observability, continuous control, and noisy observations. Therefore, the benchmark should be viewed as a focused test of structured state updates, rather than a complete test of world modelling.

Additionally, the random-uniform split is only one form of distribution shift. Uniformly random legal play produces valid trajectories far from human games, but it does not exhaust the space of possible failures. Other agents, adversarial move sequences, longer games, or targeted tests for rare rules such as castling, en passant, and promotion may reveal additional weaknesses.

Furthermore, our experiments are one use of the benchmark rather than a complete evaluation of the design space. We compare several sequence-model families under a shared supervised prediction interface, so the conclusions are controlled by a common input representation, target representation, objective, and training protocol. Broader comparisons across architectures, tokenisations, objectives, scales, and training regimes are natural next steps. CHESS-WORLD-MODEL allows these comparisons to be made in an exact and reproducible state-tracking setting.

The direct misuse and privacy risks of CHESS-WORLD-MODEL are limited, since the benchmark uses public chess games with player metadata removed and does not release a deployed decision-making system. A more plausible risk is benchmark overclaiming: strong performance on chess state reconstruction could be misrepresented as evidence of general world-modelling ability, even though chess is fully observed, deterministic, discrete, and rule-governed. The main remaining negative impact is the compute and energy cost of large-scale benchmarking.

6 Conclusion

We introduced CHESS-WORLD-MODEL, a large-scale benchmark for exact state tracking from legal chess move sequences. Models predict the complete evolving chess state at every prefix of a game, including piece placement and auxiliary variables, giving an exactly checkable test of structured latent-state maintenance.

Our experiments show that recurrent models with expressive state-transition mechanisms outperform a causal Transformer at small and medium scales under a matched interface. Held-out real-game performance saturates quickly, while the random-uniform split remains discriminative at larger scales. This shows that strong in-distribution performance can conceal failures of rule-consistent state tracking. Ablations further show that reducing recurrent transition expressivity consistently harms random-uniform performance.

These results suggest that scale can conceal poor state-tracking behaviour by allowing models to learn shortcuts that perform well in distribution without capturing the underlying transition dynamics. The random-uniform split in CHESS-WORLD-MODEL exposes these failures by testing models on legal trajectories that depart from human play while preserving the same rules. This makes the benchmark a practical setting for diagnosing state-tracking failures and for developing future sequence models with more reliable internal state updates, a necessary component of more robust world models.

References

- Assran, M., Q. Duval, I. Misra, P. Bojanowski, P. Vincent, M. Rabbat, Y. LeCun, and N. Ballas (2023). “Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Assran, M., A. Bardes, D. Fan, Q. Garrido, R. Howes, M. Komeili, M. Muckley, A. Rizvi, C. Roberts, K. Sinha, A. Zholus, S. Arnaud, A. Gejji, A. Martin, F. R. Hogan, D. Dugas, P. Bojanowski, V. Khalidov, P. Labatut, F. Massa, M. Szafraniec, K. Krishnakumar, Y. Li, X. Ma, S. Chandar, F. Meier, Y. LeCun, M. Rabbat, and N. Ballas (2025). *V-JEPA 2: Self-Supervised Video Models Enable Understanding, Prediction and Planning*. arXiv: 2506.09985.
- Cirone, N. M., A. Orvieto, B. Walker, C. Salvi, and T. Lyons (2024). “Theoretical Foundations of Deep Selective State-Space Models”. In: *Proceedings of the 38th Conference on Neural Information Processing Systems (NeurIPS)*.
- Cooper, A. (2025). *PGN2FEN: A Benchmark for Evaluating LLM Chess Reasoning*. <https://www.aidancooper.co.uk/pgn2fen-benchmark/>. Blog post and benchmark.
- Dao, T. and A. Gu (2024). “Transformers are SSMS: generalized models and efficient algorithms through structured state space duality”. In: *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Delétang, G., A. Ruoss, J. Grau-Moya, T. Genewein, L. K. Wenliang, E. Catt, C. Cundy, M. Hutter, S. Legg, J. Veness, and P. A. Ortega (2023). “Neural Networks and the Chomsky Hierarchy”. In: *Proceedings of the 11th International Conference on Learning Representations (ICLR)*.
- Edwards, S. J. (1994). *Portable Game Notation Specification and Implementation Guide*. Defines Forsyth-Edwards Notation. URL: <https://www.saremba.de/chessgml/standards/pgn/pgn-complete.htm>.
- Fan, T.-H., T.-C. Chi, and A. Rudnicky (2024). “Advancing Regular Language Reasoning in Linear Recurrent Neural Networks”. In: *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pp. 45–53.
- FIDE (2023). *FIDE Laws of Chess*. https://rcc.fide.com/wp-content/uploads/2022/11/Laws_of_Chess-2023.pdf. Effective from 1 January 2023.
- Fiekas, N. (2024). *python-chess: A Chess Library for Python*. URL: <https://github.com/niklasf/python-chess>.
- Grazzi, R., J. Siems, J. K. H. Franke, A. Zela, F. Hutter, and M. Pontil (2025). “Unlocking State-Tracking in Linear RNNs Through Negative Eigenvalues”. In: *Proceedings of the 13th International Conference on Learning Representations (ICLR)*.
- Gu, A. and T. Dao (2024). “Mamba: Linear-Time Sequence Modeling with Selective State Spaces”. In: *Proceedings of the First Conference on Language Modeling*.
- Ha, D. and J. Schmidhuber (2018). “Recurrent World Models Facilitate Policy Evolution”. In: *Proceedings of the 32nd Conference on Neural Information Processing Systems (NeurIPS)*.
- Hafner, D., T. Lillicrap, J. Ba, and M. Norouzi (2020). “Dream to Control: Learning Behaviors by Latent Imagination”. In: *Proceedings of the 8th International Conference on Learning Representations (ICLR)*.
- Harang, R., J. Naradowsky, Y. Gujju, and Y. Miyao (2025). “Tracking World States with Language Models: State-Based Evaluation Using Chess”. In: *ICML 2025 Workshop on Assessing World Models*.
- Hendrycks, D. and K. Gimpel (2016). “Gaussian Error Linear Units (GELUs)”. In: *arXiv pre-print arXiv:1606.08415*.
- Karvonen, A. (2024). “Emergent World Models and Latent Variable Estimation in Chess-Playing Language Models”. In: *Proceedings of the 1st Conference on Language Modeling (COLM)*.
- Kingma, D. P. and J. Ba (2015). “Adam: A Method for Stochastic Optimization”. In: *Proceedings of the 3rd International Conference on Learning Representations (ICLR)*.
- Lahoti, A., K. Y. Li, B. Chen, C. Wang, A. Bick, J. Z. Kolter, T. Dao, and A. Gu (2026). “Mamba-3: Improved Sequence Modeling using State Space Principles”. In: *Proceedings of the 14th International Conference on Learning Representations (ICLR)*.
- LeCun, Y. (2022). *A Path Towards Autonomous Machine Intelligence*. <https://openreview.net/forum?id=BZ5a1r-kVsf>.
- Li, B. Z., Z. C. Guo, and J. Andreas (2025). “(How) Do Language Models Track State?” In: *Proceedings of the 42nd International Conference on Machine Learning (ICML)*.

- Li, K., A. K. Hopkins, D. Bau, F. Viégas, H. Pfister, and M. Wattenberg (2023). “Emergent World Representations: Exploring a Sequence Model Trained on a Synthetic Task”. In: *Proceedings of the 11th International Conference on Learning Representations (ICLR)*.
- Lichess (2026). *Lichess Open Database*. <https://database.lichess.org/>. CC0 database exports; accessed 2026-04-20.
- Littman, M. L., R. S. Sutton, and S. Singh (2001). “Predictive Representations of State”. In: *Proceedings of the 14th Conference on Neural Information Processing Systems (NeurIPS)*.
- Loshchilov, I. and F. Hutter (2019). “Decoupled Weight Decay Regularization”. In: *Proceedings of the 7th International Conference on Learning Representations (ICLR)*.
- Merrill, W., J. Petty, and A. Sabharwal (2024). “The illusion of state in state-space models”. In: *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Merrill, W., A. Sabharwal, and N. A. Smith (2022). “Saturated Transformers are Constant-Depth Threshold Circuits”. In: *Transactions of the Association for Computational Linguistics* 10, pp. 843–856.
- Meyer-Kahlen, S. and R. Huber (2000). *Universal Chess Interface*. <https://www.shredderchess.com/chess-features/uci-universal-chess-interface.html>. UCI protocol specification.
- Nanda, N., A. Lee, and M. Wattenberg (2023). “Emergent Linear Representations in World Models of Self-Supervised Sequence Models”. In: *Proceedings of the 6th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pp. 16–30.
- Pascanu, R., T. Mikolov, and Y. Bengio (2013). “On the Difficulty of Training Recurrent Neural Networks”. In: *Proceedings of the 30th International Conference on Machine Learning (ICML)*.
- Paszke, A., S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Köpf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala (2019). “PyTorch: An Imperative Style, High-Performance Deep Learning Library”. In: *Proceedings of the 33rd Conference on Neural Information Processing Systems (NeurIPS)*.
- Shazeer, N. (2020). “GLU Variants Improve Transformer”. In: *arXiv pre-print arXiv:2002.05202*.
- Siems, J., T. Carstensen, A. Zela, F. Hutter, M. Pontil, and R. Grazi (2025). “DeltaProduct: Improving State-Tracking in Linear RNNs via Householder Products”. In: *Proceedings of the 39th Conference on Neural Information Processing Systems (NeurIPS)*.
- Srivastava, N., G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov (2014). “Dropout: A simple way to prevent neural networks from overfitting”. In: *Journal of Machine Learning Research* 15.56, pp. 1929–1958.
- Su, J., A. Murtadha, Y. Lu, S. Pan, B. Wen, and Y. Liu (2024). “RoFormer: Enhanced Transformer with Rotary Position Embedding”. In: *Neurocomputing* 568.
- Sutton, R. S. (1990). “Integrated Architectures for Learning, Planning, and Reacting Based on Approximating Dynamic Programming”. In: *Proceedings of the 7th International Conference on Machine Learning (ICML)*.
- Toshniwal, S., S. Wiseman, K. Livescu, and K. Gimpel (2022). “Chess as a Testbed for Language Model State Tracking”. In: *Proceedings of the 36th AAAI Conference on Artificial Intelligence (AAAI)*.
- Touvron, H., T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, A. Rodriguez, A. Joulin, E. Grave, and G. Lample (2023). “LLaMA: Open and Efficient Foundation Language Models”. In: *arXiv pre-print arXiv:2302.13971*.
- Vafa, K., J. Y. Chen, A. Rambachan, J. Kleinberg, and S. Mullainathan (2024). “Evaluating the World Model Implicit in a Generative Model”. In: *Proceedings of the 38th Conference on Neural Information Processing Systems (NeurIPS)*.
- Vaswani, A., N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin (2017). “Attention Is All You Need”. In: *Proceedings of the 31st Conference on Neural Information Processing Systems (NeurIPS)*.
- Walker, B., L. Yang, N. M. Cirone, C. Salvi, and T. Lyons (2025). “Structured Linear CDEs: Maximally Expressive and Parallel-in-Time Sequence Models”. In: *Proceedings of the 39th Conference on Neural Information Processing Systems (NeurIPS)*.
- Yang, S., J. Kautz, and A. Hatamizadeh (2025). “Gated Delta Networks: Improving Mamba2 with Delta Rule”. In: *Proceedings of the 13th International Conference on Learning Representations (ICLR)*.

- Yang, S. and Y. Zhang (Jan. 2024). *FLA: A Triton-Based Library for Hardware-Efficient Implementations of Linear Attention Mechanism*. URL: <https://github.com/fla-org/flash-linear-attention>.
- Zhang, B. and R. Sennrich (2019). “Root Mean Square Layer Normalization”. In: *Proceedings of the 33rd Conference on Neural Information Processing Systems (NeurIPS)*.

A Encoding details and move-token coverage

A.1 State-label layout

Table 2 gives the full target layout used by the benchmark. The target contains 64 board labels and 11 auxiliary labels, for a total of 75 categorical labels at each timestep. The represented state follows the FEN-style decomposition of a chess position, including piece placement, active colour, castling availability, en passant information, the halfmove clock, and the fullmove number (Edwards, 1994; FIDE, 2023). As in standard FEN, this representation does not encode the previous-position history needed to adjudicate repetition-based draws.

State component	Count	Vocabulary or encoding
Board squares	64	13 labels, empty or one of 12 piece types
Side to move	1	white or black
Castling rights	4	binary indicators for K, Q, k, q
En passant file	1	9 labels, none or files $a-h$
En passant rank	1	3 labels, none, rank 3, or rank 6
Halfmove clock	2	unsigned 16-bit integer
Fullmove number	2	unsigned 16-bit integer

Table 2: State representation used by the benchmark. The target includes piece placement and the auxiliary variables of the FEN-style position state, excluding the history of previous board states needed for repetition-based draws.

A.2 Move-token coverage

The input representation packs each move into a single token determined by its source square, target square, and promotion type. Because this vocabulary is intentionally overcomplete, it contains many source-target-promotion combinations that are invalid or unused in legal chess trajectories. Since the random-uniform split differs from human games, a possible concern is that it might contain packed move tokens that were never observed during training. We therefore measure identity-level move-token coverage in the training corpus and in the fixed random-uniform test set.

Let $\mathcal{V}_{\text{move}}$ denote the packed non-special move vocabulary. This vocabulary has

$$|\mathcal{V}_{\text{move}}| = 64 \times 64 \times 5 = 20,480$$

possible source-target-promotion identifiers. Table 3 reports coverage over this non-special vocabulary. The counts exclude the start and padding symbols.

Quantity	Train	Random-uniform test
Examples kept	10,000,000	10,000
Unique move identifiers	1,968/20,480	1,968/20,480
Move-token coverage	9.61%	9.61%

Table 3: Identity-level move-token coverage for the packed move vocabulary. Although the packed vocabulary contains 20,480 possible non-special move identifiers, only 1,968 appear in the training set and in the fixed random-uniform test set.

The observed move-token sets are identical across the two splits. The overlap between the training set and the random-uniform test set is

$$1,968/20,480 = 9.61\%,$$

with zero random-only move identifiers and zero train-only move identifiers. In particular, every packed move token that appears in the random-uniform test set is observed in the training data.

This rules out a simple unseen-token explanation for the random-uniform performance gap. The random-uniform split does not test zero-shot generalisation to new packed move identities. Instead, it tests whether models can apply the same observed move tokens in unfamiliar legal trajectories and board states. Thus, the random-uniform split differs from the real-game training distribution in the states and action sequences it induces, not in the set of packed move-token identities it uses.

B Additional experimental details

B.1 Data preprocessing, splits, and evaluation protocol

Each real-game example is built by replaying a PGN game from the standard initial position using the standard chess rules (Edwards, 1994; FIDE, 2023). The training and validation corpus is derived from the March 2025 Lichess open database, and the held-out real-game test set is derived from the April 2025 Lichess open database, whose public exports are released under a CC0 license (Lichess, 2026). PGN parsing, legal move generation, and board-state reconstruction are implemented using `python-chess` (Fiekas, 2024). The move sequence is packed into the categorical UCI-style move vocabulary, and every observed prefix is aligned with the chess state reached after replaying that prefix. Games shorter than 10 full moves are discarded.

The training and validation split is deterministic at the game level. Let g denote the stored game identifier. We compute

$$b(g) = \text{MD5}(g) \bmod 10,000$$

and assign a game to validation when $b(g) < 50$. This reserves approximately 0.5% of games for validation while keeping the split stable across runs. Training shuffles the order of data shards each epoch. Validation and test evaluation are deterministic.

The random-uniform test split is generated by sampling uniformly from the legal move set at each ply until termination, treating claimable draws as terminal. We then retain only games with at least 10 full moves. Final test metrics are computed on the complete held-out real-game test set of 10,000 games and the complete random-uniform test set of 10,000 games. During the first-stage sweep, validation is intentionally cheaper. Models are evaluated every 5,000 optimisation steps on 10 validation batches. With batch size 128, this corresponds to 1,280 validation games per sweep evaluation.

In addition to cross-entropy loss, we record three exactness measures. `ExactState` is the fraction of timesteps for which all 75 state labels are predicted correctly, including the 64 board-square labels and the 11 auxiliary labels. `Labelwise state accuracy` is the average fraction of the 75 state labels predicted correctly at a timestep. `Trajectory exactness` is the fraction of complete games for which `ExactState` holds at every non-padding timestep. We also record `ExactState` and labelwise state accuracy in contiguous 20-step prefix bins.

B.2 Model and optimisation details

All reported runs use AdamW (Kingma et al., 2015; Loshchilov et al., 2019) with $\beta_1 = 0.9$, global gradient clipping at 1.0 (Pascanu et al., 2013), dropout 0.1 (Srivastava et al., 2014), feed-forward multiplier 4, and batch size 128. The value of β_2 is selected by the sweep. On GPUs supporting tensor cores, TF32 matrix multiplication is enabled. Mamba-3 is trained with bfloat16 mixed precision, while the other model families are trained without mixed-precision autocasting.

The first-stage sweep consists of

$$4 \times 4 \times 4 \times 3 \times 3 \times 3 = 1728$$

one-epoch runs over model family, model size, learning rate, warmup length, β_2 , and weight decay. The learning rates are $\{10^{-4}, 3 \times 10^{-4}, 10^{-3}, 3 \times 10^{-3}\}$. The warmup lengths are $\{5000, 10000, 20000\}$. The β_2 values are $\{0.95, 0.99, 0.999\}$. The weight decays are $\{0, 10^{-3}, 10^{-2}\}$. These runs use linear warmup followed by an effectively constant learning rate. The sweep uses a fixed seed to make architecture comparisons less sensitive to run-to-run variation.

In the second stage, selected first-stage configurations are retrained from scratch for four epochs using one of two cosine-decay schedules, with final learning rates equal to 0.1 or 0.01 times the maximum learning rate. Each run terminates at optimisation step 310,944. The final reported configuration for each family and size is selected by held-out validation loss among the completed second-stage runs.

Transformer. The Transformer baseline is a small LLaMA-like causal decoder for autoregressive world-state prediction (Vaswani et al., 2017; Touvron et al., 2023). It uses causal self-attention, RoPE positional embeddings (Su et al., 2024), RMSNorm (Zhang et al., 2019), pre-norm residual blocks, SwiGLU MLPs (Shazeer, 2020), and a final RMSNorm before the structured state-prediction heads. The four sizes use

$$(d, n_{\text{layers}}, n_{\text{heads}}) = (128, 1, 2), (256, 2, 4), (384, 4, 6), (512, 6, 8).$$

Model	d128/l1	d256/l2	d384/l4	d512/l6
Transformer	3.13	7.83	18.05	36.66
SLiCE	3.14	7.96	18.64	38.24
Mamba-3	3.20	8.25	19.72	40.83
Gated DeltaNet	3.16	7.70	16.89	35.14
Diagonal SLiCE	3.03	7.04	14.50	27.21
Mamba-2	3.20	8.23	19.64	40.65
Gated DeltaNet with positive eigenvalues	3.16	7.70	16.89	35.14

Table 4: Exact parameter counts in millions for the configurations used in the paper plots. The Gated DeltaNet sign ablation leaves the parameter count unchanged, Mamba-2 is nearly parameter matched to Mamba-3, and diagonal SLiCE is smaller than block-diagonal SLiCE because its transition parameterisation contains fewer parameters.

SLiCE. The main SLiCE model uses input-dependent block-diagonal recurrent transitions (Walker et al., 2025). We use block size 8, and evaluate the recurrent computation in parallel chunks of length 256. Each layer uses a GELU feed-forward MLP (Hendrycks et al., 2016), RMSNorm (Zhang et al., 2019), and dropout on the residual path (Srivastava et al., 2014). The diagonal SLiCE ablation replaces the block-diagonal transition with a diagonal transition while keeping the rest of the configuration fixed.

Gated DeltaNet. The main Gated DeltaNet model uses the same head schedule as the Transformer, together with head dimension 48, value expansion factor 2, gating, a short convolution of width 4, and chunked recurrent computation (Yang et al., 2025). The main variant allows negative transition eigenvalues, following recent evidence that signed eigenvalue structure improves state tracking in linear recurrent models (Grazzi et al., 2025). The ablation restricts the transition eigenvalues to be positive while leaving the parameter count unchanged.

Mamba. The main Mamba model uses Mamba-3 with the MIMO variant and rank 4 (Lahoti et al., 2026). It uses state dimension 128, convolution width 4, expansion factor 2, head dimension 64, one group, intermediate dimension $4d$, RMSNorm (Zhang et al., 2019), and FP32 residual accumulation. The Mamba-2 ablation keeps the same nominal size schedule and replaces the Mamba-3 block with Mamba-2 (Gu et al., 2024; Dao et al., 2024).

Table 4 provides the exact parameter count for each model considered in this paper.

All models were trained on a single GPU. Training used multiple GPU types, depending on model size and kernel requirements. The most restrictive runs were the Mamba-3 configurations, since the MIMO kernel requires an NVIDIA H100 GPU. The longest second-stage retrains took approximately four days. The experiments use the following publicly available datasets, libraries, and baseline implementations:

- Lichess Open Database (Lichess, 2026). License: CC0.
URL: <https://database.lichess.org/>
- python-chess (Fiekas, 2024). License: GPL-3.0-or-later.
URL: <https://github.com/niklasf/python-chess>
- torch-slices (Walker et al., 2025). License: MIT.
URL: <https://github.com/datasig-ac-uk/slices>
- mamba-ssm (Gu et al., 2024; Dao et al., 2024). License: Apache 2.0.
URL: <https://github.com/state-spaces/mamba>
- flash-linear-attention (Yang et al., 2024). License: MIT.
URL: <https://github.com/fla-org/flash-linear-attention>
- pytorch (Paszke et al., 2019). License: BSD-style.
URL: <https://github.com/pytorch/pytorch>

Model	Size	Val. loss	ExactState (%)	Labelwise (%)	Trajectory (%)	LR	Warmup	β_2	Weight decay
Transformer	d128/l1	0.143939	26.02	95.61	0.16	0.003	5000	0.95	0
Transformer	d256/l2	0.023679	67.15	99.22	10.62	0.003	5000	0.95	0.001
Transformer	d384/l4	0.000959	98.15	99.97	77.19	0.003	5000	0.999	0.001
Transformer	d512/l6	0.000237	99.51	99.99	92.19	0.001	5000	0.95	0
SLiCE	d128/l1	0.061826	38.65	98.03	0.70	0.001	20000	0.999	0.01
SLiCE	d256/l2	0.004505	90.71	99.85	34.69	0.003	10000	0.99	0.001
SLiCE	d384/l4	0.000760	98.45	99.98	78.44	0.001	5000	0.95	0.001
SLiCE	d512/l6	0.000359	99.36	99.99	88.59	0.001	10000	0.95	0.001
Mamba-3	d128/l1	0.047153	48.92	98.43	2.81	0.003	10000	0.95	0
Mamba-3	d256/l2	0.003149	93.58	99.90	48.28	0.003	5000	0.999	0
Mamba-3	d384/l4	0.000408	99.14	99.99	85.55	0.001	5000	0.95	0.01
Mamba-3	d512/l6	0.000143	99.70	100.00	92.81	0.003	5000	0.95	0.001
Gated DeltaNet	d128/l1	0.061575	45.51	98.01	2.42	0.001	10000	0.99	0
Gated DeltaNet	d256/l2	0.001908	95.98	99.94	57.03	0.003	20000	0.99	0
Gated DeltaNet	d384/l4	0.000448	99.01	99.99	84.69	0.001	5000	0.95	0
Gated DeltaNet	d512/l6	0.000204	99.54	99.99	91.56	0.001	10000	0.95	0

Table 5: Best first-stage sweep configuration in each family and size cell, ranked by validation loss on the held-out real-game split. ExactState is the strict timestep metric requiring all 75 state labels to be correct. Labelwise state accuracy and trajectory exactness are measured on the same 1280-game validation snapshot used during the sweep.

C Hyperparameter sweep and full run results

Table 5 reports the best one-epoch configuration for every family and size when ranked by validation loss. The broad pattern matches the main text. Recurrent models dominate the Transformer at small scale, and all models perform well on the in-distribution validation sweep at large scales.

Validation-loss ranking does not coincide exactly with ranking by ExactState or labelwise state accuracy. Across the 16 family-size cells, the validation-loss winner is also the ExactState winner in 5 cases and the labelwise winner in 9 cases. In practice, the disagreements are between nearby configurations within the same family and size cell. We therefore use validation loss as the single model-selection criterion.

The second stage retrains selected first-stage configurations for four epochs. For every family and size except SLiCE d512/l6, the selected configuration is the first-stage validation-loss winner in Table 5. For SLiCE d512/l6, the first-stage validation-loss winner was unstable during the full four epoch run, so we instead chose the second-best first-stage configuration by validation loss. For each selected configuration, we compare two cosine-decay floors, equal to 0.1 and 0.01 times the maximum learning rate. Table 6 compares the two second-stage cosine-decay floors for each selected first-stage configuration. The preferred floor depends on family and scale. The final reported baseline in each family and size cell is the completed full run with lower held-out validation loss.

Table 7 consolidates the final optimisation hyperparameters used by the reported baseline models.

D Complete test-set results

Table 8 reports the complete test results for the main model families. In all four families, the held-out real-game split becomes nearly saturated by d384/l4, while the random-uniform split remains more discriminative. This is most visible in trajectory exactness, which remains low on the random-uniform split even when timestep-level held-out real-game metrics are close to perfect.

E Variant results

Tables 9 and 10 give the full ablation results for the less expressive recurrent variants. These ablations reduce transition expressivity while keeping the rest of the training and evaluation protocol fixed. The random-uniform split is more sensitive to these changes than the held-out real-game split. For SLiCE and Mamba, the less expressive variant is substantially worse across scales. For Gated DeltaNet, the positive-eigenvalue variant is worse in loss and ExactState through d384/l4 and is nearly tied at d512/l6.

Model	Size	Best floor	Best val. loss	Other floor	Other val. loss	ExactState (%)	Labelwise (%)
Transformer	d128/l1	0.1	0.136054	0.01	0.136446	26.735	95.8573
Transformer	d256/l2	0.01	0.015992	0.1	0.016065	74.984	99.4716
Transformer	d384/l4	0.01	5.60×10^{-5}	0.1	1.14×10^{-4}	99.876	99.9982
Transformer	d512/l6	0.1	7.47×10^{-6}	0.01	7.64×10^{-6}	99.980	99.9997
SLiCE	d128/l1	0.1	0.054437	0.01	0.055932	41.269	98.2439
SLiCE	d256/l2	0.01	0.001462	0.1	0.001529	96.726	99.9493
SLiCE	d384/l4	0.1	6.98×10^{-5}	0.01	7.28×10^{-5}	99.841	99.9977
SLiCE	d512/l6	0.01	1.32×10^{-5}	0.1	1.42×10^{-5}	99.971	99.9996
Mamba-3	d128/l1	0.01	0.041538	0.1	0.042556	52.194	98.6139
Mamba-3	d256/l2	0.1	0.000797	0.01	0.000804	98.216	99.9730
Mamba-3	d384/l4	0.01	1.18×10^{-5}	0.1	2.69×10^{-5}	99.982	99.9997
Mamba-3	d512/l6	0.1	7.27×10^{-7}	0.01	8.27×10^{-6}	99.999	100.0000
Gated DeltaNet	d128/l1	0.01	0.050420	0.1	0.050950	49.955	98.3648
Gated DeltaNet	d256/l2	0.01	0.000375	0.1	0.000539	99.207	99.9888
Gated DeltaNet	d384/l4	0.1	9.61×10^{-6}	0.01	1.42×10^{-5}	99.977	99.9997
Gated DeltaNet	d512/l6	0.1	1.81×10^{-6}	0.01	2.43×10^{-6}	99.998	100.0000

Table 6: Second-stage full run comparison. Each row compares the two cosine-decay floors used in the full four-epoch runs for the selected first-stage configuration. ExactState is the strict timestep metric requiring all 75 state labels to be correct. The reported ExactState and labelwise values are the validation metrics of the better full runs.

Model	Size	LR	Warmup	β_2	Weight decay	LR floor
Transformer	d128/l1	0.003	5000	0.95	0	0.1
Transformer	d256/l2	0.003	5000	0.95	0.001	0.01
Transformer	d384/l4	0.003	5000	0.999	0.001	0.01
Transformer	d512/l6	0.001	5000	0.95	0	0.1
SLiCE	d128/l1	0.001	20000	0.999	0.01	0.1
SLiCE	d256/l2	0.003	10000	0.99	0.001	0.01
SLiCE	d384/l4	0.001	5000	0.95	0.001	0.1
SLiCE	d512/l6	0.0003	5000	0.99	0.001	0.01
Mamba-3	d128/l1	0.003	10000	0.95	0	0.01
Mamba-3	d256/l2	0.003	5000	0.999	0	0.1
Mamba-3	d384/l4	0.001	5000	0.95	0.01	0.01
Mamba-3	d512/l6	0.003	5000	0.95	0.001	0.1
Gated DeltaNet	d128/l1	0.001	10000	0.99	0	0.01
Gated DeltaNet	d256/l2	0.003	20000	0.99	0	0.01
Gated DeltaNet	d384/l4	0.001	5000	0.95	0	0.1
Gated DeltaNet	d512/l6	0.001	10000	0.95	0	0.1

Table 7: Final selected optimisation hyperparameters for the reported baseline models. The LR floor is the final learning rate as a fraction of the maximum learning rate in the second-stage cosine schedule.

Model	Size	Held-out real games				Random-uniform games			
		Loss	Labelwise (%)	ExactState (%)	Trajectory (%)	Loss	Labelwise (%)	ExactState (%)	Trajectory (%)
Transformer	d128/l1	0.137515	95.81	26.82	0.02	0.611882	85.53	2.67	0.00
Transformer	d256/l2	0.016659	99.45	74.44	11.91	0.280601	92.48	8.58	0.06
Transformer	d384/l4	7.65×10^{-5}	100.00	99.86	97.35	0.036100	99.05	59.45	1.79
Transformer	d512/l6	2.93×10^{-5}	100.00	99.96	99.61	0.029630	99.33	72.93	4.87
SLiCE	d128/l1	0.055239	98.22	41.20	0.16	0.322331	92.21	4.14	0.00
SLiCE	d256/l2	0.001636	99.95	96.57	63.95	0.075425	98.19	33.16	0.52
SLiCE	d384/l4	8.40×10^{-5}	100.00	99.85	97.27	0.035920	99.45	70.85	2.54
SLiCE	d512/l6	2.67×10^{-5}	100.00	99.95	98.97	0.032891	99.49	73.47	4.37
Mamba-3	d128/l1	0.042266	98.57	51.96	1.48	0.305990	91.57	5.66	0.00
Mamba-3	d256/l2	0.000908	99.97	98.13	77.35	0.068237	98.10	33.31	0.90
Mamba-3	d384/l4	3.66×10^{-5}	100.00	99.94	99.04	0.029089	99.30	68.78	5.07
Mamba-3	d512/l6	1.29×10^{-5}	100.00	99.99	99.76	0.015915	99.66	81.13	10.82
Gated DeltaNet	d128/l1	0.051443	98.34	49.47	1.04	0.534119	90.37	5.05	0.00
Gated DeltaNet	d256/l2	0.000467	99.99	99.12	87.46	0.086148	97.90	36.15	1.18
Gated DeltaNet	d384/l4	5.79×10^{-5}	100.00	99.93	99.46	0.044545	99.04	56.87	4.55
Gated DeltaNet	d512/l6	3.43×10^{-5}	100.00	99.96	99.60	0.035822	99.22	61.19	5.71

Table 8: Complete test results for the main model families. Loss is the average cross-entropy over the 75 state-label heads. Labelwise accuracy is the average percentage of state labels recovered at a timestep. ExactState is the strict timestep metric requiring all 75 state labels to be correct. Trajectory exactness is the percentage of games for which ExactState holds at every non-padding timestep.

Variant	Size	Loss	Labelwise (%)	ExactState (%)	Trajectory (%)
Diagonal SLiCE	d128/l1	0.120725	96.48	16.67	0.00
Diagonal SLiCE	d256/l2	0.042201	99.76	86.40	24.99
Diagonal SLiCE	d384/l4	0.000582	99.98	98.89	86.67
Diagonal SLiCE	d512/l6	0.000457	99.99	99.33	92.69
Mamba-2	d128/l1	0.056944	98.13	46.41	1.04
Mamba-2	d256/l2	0.002905	99.90	94.40	53.55
Mamba-2	d384/l4	0.000149	100.00	99.80	97.54
Mamba-2	d512/l6	9.05×10^{-5}	100.00	99.88	98.63
Gated DeltaNet with positive eigenvalues	d128/l1	0.053603	98.26	47.49	0.78
Gated DeltaNet with positive eigenvalues	d256/l2	0.001104	99.96	97.87	76.00
Gated DeltaNet with positive eigenvalues	d384/l4	0.000110	100.00	99.85	98.03
Gated DeltaNet with positive eigenvalues	d512/l6	3.58×10^{-5}	100.00	99.95	99.56

Table 9: Held-out real-game test results for the less expressive recurrent variants. ExactState is the strict timestep metric requiring all 75 state labels to be correct.

Variant	Size	Loss	Labelwise (%)	ExactState (%)	Trajectory (%)
Diagonal SLiCE	d128/l1	35.9174	87.33	2.09	0.00
Diagonal SLiCE	d256/l2	4.43×10^{17}	88.96	13.32	0.13
Diagonal SLiCE	d384/l4	8.23×10^6	97.76	38.86	1.17
Diagonal SLiCE	d512/l6	1.615620	98.33	39.14	1.80
Mamba-2	d128/l1	0.393216	89.85	4.92	0.00
Mamba-2	d256/l2	0.125556	96.41	19.32	0.44
Mamba-2	d384/l4	0.057472	98.60	45.66	2.48
Mamba-2	d512/l6	0.052418	98.86	50.02	3.04
Gated DeltaNet with positive eigenvalues	d128/l1	0.496955	90.39	4.80	0.00
Gated DeltaNet with positive eigenvalues	d256/l2	0.116137	97.25	27.79	0.70
Gated DeltaNet with positive eigenvalues	d384/l4	0.060025	98.61	48.57	2.93
Gated DeltaNet with positive eigenvalues	d512/l6	0.036353	99.22	61.17	6.28

Table 10: Random-uniform test results for the less expressive recurrent variants. ExactState is the strict timestep metric requiring all 75 state labels to be correct. The diagonal SLiCE ablation is notably unstable out of distribution. Its ExactState drops sharply, and its random-uniform cross-entropy can become very large even when held-out real-game performance remains strong.

Model	0–20	80–100	160–180	240–260	320–340	400–420	500–520
Transformer	99.99	89.75	79.95	72.82	37.61	0.04	0.00
SLiCE	99.98	86.50	80.17	82.01	34.53	14.34	0.00
Mamba-3	100.00	96.01	88.99	88.26	46.44	20.74	0.00
Gated DeltaNet	99.98	95.07	86.80	13.23	1.25	0.00	0.00

Table 11: Random-uniform d512 ExactState in selected 20-step prefix bins. Values are percentages. ExactState requires all 75 state labels to be correct.

Model	0–20	80–100	160–180	240–260	320–340	400–420	500–520
Transformer	100.00	99.85	99.68	99.56	98.69	96.04	93.48
SLiCE	100.00	99.79	99.67	99.67	98.79	98.06	96.18
Mamba-3	100.00	99.94	99.83	99.80	99.05	98.44	97.23
Gated DeltaNet	100.00	99.93	99.80	98.60	97.86	97.29	96.08

Table 12: Random-uniform d512 labelwise state accuracy in the same selected 20-step prefix bins as Table 11. Values are percentages.

F Extended temporal-bin metrics

The timestep-averaged metrics in Table 8 hide where errors first appear along a trajectory. Tables 11 and 12 therefore report saved 20-step prefix-bin metrics on the random-uniform test set for the four d512 baseline models.

Two features stand out. First, ExactState can collapse long before labelwise accuracy does, which means that many long-prefix errors are caused by a small number of wrong state labels. Second, different architectures fail at different horizons. Gated DeltaNet remains strong through roughly 180 plies and then degrades abruptly, whereas Mamba-3 and SLiCE decay more gradually.